



Machine Learning for Scientists and Engineers

August 2024

Name	ID	Email
João Felipe Gueiros	941200164	joao.g@campus.technion.ac.il

1 Theory

1st Problem

Problem 2.1 To walk “downhill” on the loss function (equation 2.5), we measure its gradient with respect to the parameters ϕ_0 and ϕ_1 . Calculate expressions for the slopes $\partial L/\partial\phi_0$ and $\partial L/\partial\phi_1$.

Solution:

$$\frac{\partial L}{\partial\phi_0} = \frac{\partial \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2}{\partial\phi_0} = \frac{\sum_{i=1}^I \partial(\phi_0 + \phi_1 x_i - y_i)^2}{\partial\phi_0} = \sum_{i=1}^I (2\phi_0 + 2\phi_1 x_i - 2y_i)$$

$$\frac{\partial L}{\partial\phi_1} = \frac{\partial \sum_{i=1}^I (\phi_0 + \phi_1 x_i - y_i)^2}{\partial\phi_1} = \frac{\sum_{i=1}^I \partial(\phi_0 + \phi_1 x_i - y_i)^2}{\partial\phi_1} = \sum_{i=1}^I (2\phi_0 x_i + 2\phi_1 x_i^2 - 2y_i x_i)$$

2nd Problem

Problem 2.2 Show that we can find the minimum of the loss function in closed form by setting the expression for the derivatives from problem 2.1 to zero and solving for ϕ_0 and ϕ_1 . Note that this works for linear regression but not for more complex models; this is why we use iterative model fitting methods like gradient descent (figure 2.4).

Solution:

Let's set equate the expressions to 0:

$$\frac{\partial L}{\partial\phi_0} = \sum_{i=1}^I (2\phi_0 + 2\phi_1 x_i - 2y_i) = 0 \quad (1)$$

$$\frac{\partial L}{\partial\phi_1} = \sum_{i=1}^I (2\phi_0 x_i + 2\phi_1 x_i^2 - 2y_i x_i) = 0 \quad (2)$$

Remark: I'll just set the $\sum_{i=1}^I$ to \sum to shorten my typing.

Back to equation (1)

$$\begin{aligned} \sum (y_i) &= \sum \phi_0 + \sum \phi_1 x_i \Rightarrow \sum \phi_0 = \sum y_i - \phi_1 \sum x_i \Rightarrow I \cdot \phi_0 = \sum y_i - \phi_1 \sum x_i \\ \phi_0 &= \frac{\sum y_i - \phi_1 \sum x_i}{I} \end{aligned}$$

Now, we must find ϕ_1

Reorganizing equation (2), we get:

$$\phi_1 = \frac{\sum y_i x_i - \phi_0 \sum x_i}{\sum x_i^2}$$

Substituting ϕ_0 from before we get:

$$\phi_1 = \frac{\sum y_i x_i - \frac{\sum y_i - \phi_0}{I} \sum x_i}{\sum x_i^2}$$

Reorganizing for ϕ_1

$$\phi_1 = \frac{I \sum y_i x_i - \sum x_i \sum y_i}{I \sum x_i^2 - (\sum x_i)^2}$$

□

2 Computation

Question 1.

```

HW2 > hw_2_question_1.py > ...
1  import numpy as np
2  import matplotlib.pyplot as plt
3  # Getting data
4  x = np.array([1, 2, 3, 4, 5, 6])
5  y = np.array([1, 3, 2, 5, 4, 6])
6  data = np.concatenate((x[np.newaxis, :], y[np.newaxis, :]), axis=0)
7  ## Calculating stuff
8  x_avg = np.mean(x)
9  y_avg = np.mean(y)
10 beta_1_num = np.sum((x - x_avg) * (y - y_avg))
11 beta_1_den = np.sum((x - x_avg) ** 2)
12 beta_1 = beta_1_num / beta_1_den
13 beta_0 = y_avg - beta_1*x_avg
14 x_line = np.linspace(1,6,100)
15
16 y_line = beta_1*x_line + beta_0
17 plt.plot(x_line,y_line,"-r",label = "Best line")
18 plt.plot()
19 plt.scatter(x, y, color='blue', label='Points')
20 plt.grid()
21 plt.xlabel('x')
22 plt.ylabel('y')
23 plt.title('Linear regression for blue points')
24 plt.show()

```

Figure 1: Python program to implement linear regression

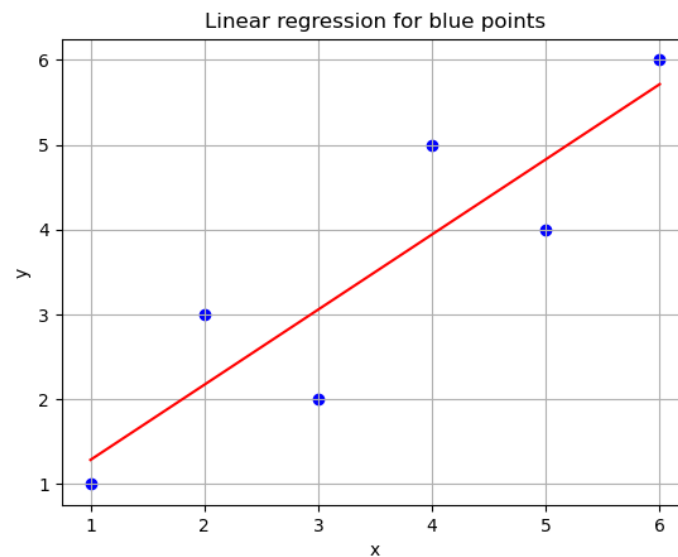


Figure 2: Plot of linear regression

Question 2.

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 x = np.array([1, 2, 3, 4, 5, 6])
4 y = np.array([1, 3, 2, 5, 4, 6])
5
6 def cost_function(beta_0, beta_1, x, y):
7     h = beta_0 * 1 + beta_1 * x
8     J = 1 / (2 * len(x)) * np.sum((h - y) ** 2)
9
10    return J
11
12 def gradient_descent(alpha, iter, beta_0, beta_1, x, y, x0):
13     cost = np.zeros(iter)
14     for i in range(iter):
15         h = beta_0 * 1 + beta_1 * x
16         gradient_0 = 1 / (len(x)) * np.sum(h - y) * x0
17         beta_0 = beta_0 - alpha * gradient_0
18         gradient_1 = 1 / (len(x)) * np.sum((h - y) * x)
19         beta_1 = beta_1 - alpha * gradient_1
20         cost[i] = cost_function(beta_0, beta_1, x, y)
21
22     return [cost, beta_0, beta_1]
23
24 alpha = 0.01
25 iter = 1000
26 lins = np.linspace(1, iter, iter)
27 x0 = 1
28 beta_0 = 1
29 beta_1 = 1
30 [cost, beta_0, beta_1] = gradient_descent(alpha, iter, beta_0, beta_1, x, y, x0)
31 x_line = np.linspace(1, 6, 100)
32 y_line = beta_1 * x_line + beta_0
33
34 # Figure 1: Plot of cost
35 plt.figure(1) # Start a new figure
36 plt.plot(lins, cost, label="Cost")
37 plt.grid()
38 plt.xlabel('Iterations')
39 plt.ylabel('Cost')
40 plt.title('Cost Convergence')
41 plt.legend()
42 plt.show()
43
44 # Figure 2: Linear regression fit
45 plt.figure(2) # Start a new figure
46 plt.plot(x_line, y_line, "-r", label="Best Line") # Best-fit line
47 plt.scatter(x, y, color="blue", label="Data Points") # Data points
48 plt.grid()
49 plt.xlabel("x")
50 plt.ylabel("y")
51 plt.title("Linear Regression for Blue Points")
52 plt.legend()
53 plt.show()

```

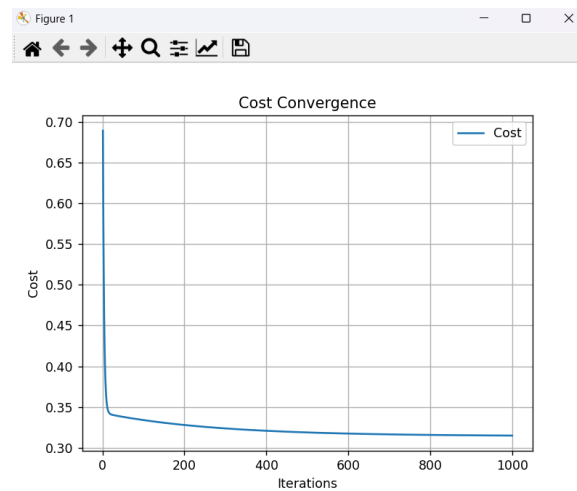
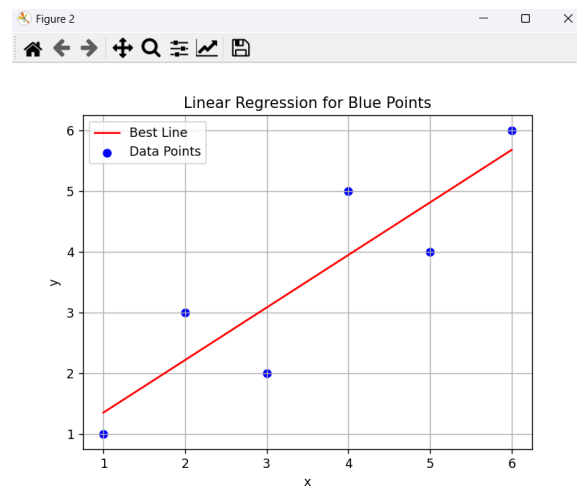
Figure 3: Cost convergence plot with $\alpha = 0.01$ 

Figure 4: Linear regression with gradient descent plot

Question 4.

```

1 import torch
2 import torch.nn as nn
3
4 # Define neural network
5 class Question4(nn.Module):
6     def __init__(self, in_dim=2, HL1_dim=2, out_dim=1):
7         super().__init__()
8         self.output = nn.Sequential(
9             nn.Linear(in_dim, HL1_dim), nn.Sigmoid(),
10            nn.Linear(HL1_dim, out_dim), nn.Sigmoid()
11        )
12

```

```

13     # Initialize weights and biases for the first layer
14     self.output[0].weight.data = torch.tensor([[0.1, 0.3], [0.2, 0.4]]) # W1
15     self.output[0].bias.data = torch.tensor([0.5, 0.6]) # B1
16
17     # Initialize weights and biases for the second layer
18     self.output[2].weight.data = torch.tensor([[0.7, 0.8]]) # W2
19     self.output[2].bias.data = torch.tensor([0.9]) # B2
20
21     def forward(self, x):
22         return self.output(x)
23
24 # Input tensor
25 x = torch.tensor([[0.3, 0.7]])
26
27 # Target value
28 target = torch.tensor([[1.0]]) # Expected output
29
30 # Define model
31 model = Question4()
32
33 # Perform forward pass
34 output = model(x)
35
36 # Compute error using the given formula
37 error = 0.5 * ((output - target) ** 2).sum() # Squared error
38
39 # Print output and error
40 print(f"Output: {output.item()}")
41 print(f"Error: {error.item()}")

```

```

PS C:\Users\Joao Felipe\OneDrive - Technion\Documents\Winter 2024 - 2025\Machine Learning> & C:/anaconda3/python.exe "c:/Users/Joao Felipe/OneDrive - Technion/Docu
ments/Winter 2024 - 2025/Machine Learning/HW2/hw_2_question_4.py"
Output: 0.8753568530082703
Error: 0.007767957169562578
PS C:\Users\Joao Felipe\OneDrive - Technion\Documents\Winter 2024 - 2025\Machine Learning>

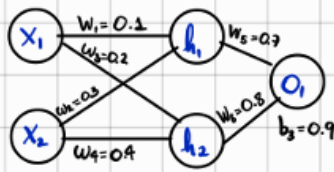
```

Figure 5: Result of error

I obtained a very similar result to the hand calculation, probably this slight difference is due to my approximations in the hand calculation.

Question 3.

Written solution in the following 2 pages



$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T = \begin{bmatrix} 0.3 & 0.7 \end{bmatrix}$$

$$W_1 = \begin{bmatrix} w_1 & w_2 \\ w_3 & w_4 \end{bmatrix} = \begin{bmatrix} 0.1 & 0.2 \\ 0.3 & 0.4 \end{bmatrix}$$

$$B_1 = [b_1 \ b_2] = [0.5 \ 0.6]$$

$$W_2 = \begin{bmatrix} w_5 \\ w_6 \end{bmatrix} = \begin{bmatrix} 0.7 \\ 0.8 \end{bmatrix}$$

$$B_2 = [b_3] = [0.9]$$

Θ_{ij} , i is the # of the hidden unit and j is # of the input

h_i , i is the # of the hidden unit

$$h_d = a \left[\Theta_{d0} + \sum_{i=1}^{D_i (\# \text{ inputs})} \Theta_{di} x_i \right], \quad y_j = \phi_{j0} + \sum_{d=1}^D \phi_{jd} h_d$$

$$h_1 = a [\Theta_{10} + \Theta_{11} x_1 + \Theta_{12} x_2]$$

$$h_2 = a [\Theta_{20} + \Theta_{21} x_1 + \Theta_{22} x_2]$$

$$y_o = \phi_{10} + \phi_{o1} h_1 + \phi_{o2} h_2$$

Figure 6: Question 3 solution part 1

$$\text{So, } h_1 = a[b_1 + w_1 x_1 + w_2 x_2]$$

$$h_2 = a[b_2 + w_3 x_1 + w_4 x_2]$$

$$\Rightarrow h_1 = a[0.5 + 0.1 \cdot 0.3 + 0.3 \cdot 0.7] = a[0.74]$$

$$h_2 = a[0.6 + 0.2 \cdot 0.3 + 0.4 \cdot 0.7] = a[0.94]$$

$$\Rightarrow h_1 = \frac{1}{1 + e^{-0.74}} \approx 0.677$$

$$h_2 = \frac{1}{1 + e^{-0.94}} \approx 0.719$$

So,

$$O_0 = a[y_0] = a[\phi_0 + \phi_1 h_1 + \phi_2 h_2]$$

$$\Rightarrow y_0 = b_3 + w_5 h_1 + w_6 h_2 = 0.9 + 0.7 \cdot (0.677) + 0.8(0.719)$$

$$\Rightarrow y_0 = 1.941$$

$$O_0 = a[1.941] = \frac{1}{1 + e^{-1.941}} \approx 0.874$$

$$\text{error} = \frac{1}{2} (0.874 - 1)^2 \approx 0.00788 //$$

Figure 7: Question 3 solution part 2